
What Survives the Filter? Trivial Baselines on Stored Honest/Attack Pairs in AI-Control Monitoring

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Trusted monitoring in AI control is evaluated on stored honest/attack pairs. The
2 field already knows these pairs leak: it filters to working solutions and strips
3 comments. We ask what that preprocessing leaves behind. On the field’s own
4 filtered, comment-stripped split, a *trained* bag-of-words classifier scores above
5 roughly four-fifths of the open-weight monitors ranked on it, and a classifier
6 that reads no code at all — only problem metadata — scores above most of
7 them. We trace the residue to authorship rather than sabotage: the honest class
8 is human-written while every attack is model-written. The signal transfers across
9 backdoor generators, is nearly as strong on inert backdoors as on working ones,
10 and survives the field’s comment stripper. Regenerating only the honest half does
11 not isolate content: with an open-weight model writing the honest class against the
12 stored backdoors, separability *risks*. Holding authorship constant by construction,
13 with one model writing both classes, removes between a third (Claude Sonnet 5,
14 $\rho = 0.687$) and nine-tenths (Claude Haiku 4.5, $\rho = 0.095$) of the separability;
15 matching the attack prompt’s framing removes most of what remains, and what
16 survives on attacks that actually fire is surface content a monitor also reads. A zero-
17 shot trusted monitor is inflated the same way, scoring 0.9579 on the stored pairs and
18 0.818 on firing backdoors from the same author. We recommend that stored-pair
19 evaluations report partial-input and bag-of-words floors against permutation nulls,
20 use same-generator honest samples and the same prompt, and drop inert backdoors
21 from both classes.

22 1 Introduction

23 Trusted monitoring is the backbone protocol of AI control: a weaker trusted model reads an untrusted
24 model’s output and flags sabotage [4]. It is developed, tuned and ranked almost entirely on *stored*
25 honest/attack pairs — fixed corpora of solutions, some carrying a planted backdoor [12, 19, 9]. The
26 field already knows these pairs leak. Koran et al. keep only working honest solutions and working
27 backdoors and strip every comment, with the rationale stated [12]; Control Tax reports the pass-tests
28 pre-filter and appendices “Monitoring failed backdoors” [19]; [18] is an entire paper on the comment
29 channel. The sharpest treatment is Jha’s, whose harness ships a `-leak-check` flag with two trivial
30 baselines inside it and concludes that stripping is mandatory rather than cosmetic: a `#-count` detector
31 falls from 0.6957 to 0.4983 once comments go, and code length alone sits at 0.4726 [9]. This paper
32 starts from that work.

33 What nobody has measured is what the preprocessing leaves behind. On the field’s own filtered,
34 comment-stripped split, a bag-of-words classifier reaches AUROC 0.806 — above 20 of the 24
35 open-weight monitors ranked on that same split, whose median is 0.609 — and a classifier that reads
36 no code at all, only problem difficulty, test count and question length, reaches 0.772, above 20.




Figure 1: **Two floors against the 24 open-weight monitors ranked on the same split.** AUROC on Koran et al.’s 1,375-item test split (172 attack / 1,203 honest), comments stripped; monitor scores from [9] (MIT), median 0.609, max 0.937. **Two caveats.** (i) The monitors are zero-shot prompted while the bag-of-words classifier is *trained* on the benchmark’s own paired train split: a statement about what the corpus makes recoverable [6], not a claim that bag-of-words is a better monitor. (ii) The no-code floor is cross-validated *within* the test split (appendix A); on one held-out slice it is 0.716 (CI 0.6476–0.7796; $p = 0.030$), above 18 of 24, 12 with a paired interval excluding zero.

37 Much of the residue is authorship, not sabotage. Honest solutions come from APPS [7] and are
 38 human-written; every backdoor is model-written. The classifier transfers across backdoor generators
 39 and is nearly as strong on **inert** backdoors as on working ones — neither is a property of a backdoor.
 40 That bag-of-words separates human-written from model-written code is established out-of-field
 41 [14]; the known-ness of the mechanism is the argument, not a dent in it. The neighbouring field
 42 already applies the control missing here — CodeMirage gates its classes on line and character count
 43 [5], MultiAIGCD pairs to shared problem definitions [1] — so unmatched pairs are a defect by a
 44 published standard. Adjacent control work attacks other axes: [20] subverts monitors with injected
 45 text, [3] shows attacker-side selection inflates safety estimates, and [15] removes transcript length as
 46 a nuisance covariate; we publish one as a per-benchmark floor.

47 We contribute (i) a protocol $\times$ baseline measurement of what survives the field’s preprocessing;
 48 (ii) a same-generator control on two generators, attributing between a third and nine-tenths of the
 49 separability to authorship and most of the remainder to the attack prompt’s framing, with a negative
 50 result on the way: regenerating one half of a stored pair swaps the confound rather than removing
 51 it; (iii) a second selection artifact the filter installs by defining the negative class through attack
 52 failure; and (iv) a preprocessing checklist and a released split. **Scope:** this concerns monitor training,
 53 selection and benchmarking on *stored* pairs; our measurement is of one APPS-derived family plus
 54 five agentic settings. ControlArena’s live evaluations generate both classes and are not implicated.

55 2 Setup

56 **The confound is legible in the schema.** The canonical stored-pair artifact for APPS-based control
 57 evaluations [8] records `model` and `source` on every backdoor and no generator field on any honest
 58 solution (Table 1, appendix).

59 **Artifacts and provenance.** We use seven public corpora; their per-class provenance is the object of
 60 study: in the APPS family the honest class is human and the attack class is one of ten named models.
 61 **Six of the seven, from two publishers, declare no licence** [11, 17]; we report only aggregates over
 62 them and redistribute no corpus.

63 **Evaluation surface, and naming protocols.** Our headline protocol is the filtered, comment-
 64 stripped split of [12]: 1,209 train problems (2,418 paired items) and 1,375 test items (172 attack,
 65 1,203 honest), labelled by `backdoor_works`, reconstructed with the released code of [9]. Every
 66 comparison below names its split, metric and preprocessing. Generated code is scored by our own
 67 runner, applied to both classes (appendix A).

68 3 Results

69 Table 2 (appendix) reports what a trivial classifier recovers under each protocol. The fix works, on
70 the two carriers §1 names. What it does not reach, on that same split, is a bag-of-words floor of
71 0.806 (95% CI 0.7704–0.8409 by bootstrap over problems; permutation null 95th percentile 0.5465,
72 $p < 0.005$), above 20 of the 24 monitors by point estimate and 18 of them with a paired-bootstrap
73 interval excluding zero, and a no-code floor of 0.772 (CI 0.7319–0.8005 on the pooled out-of-fold
74 0.766; null 0.5545, $p < 0.005$), above 20 and 17 respectively (Figure 1): the preprocessing removed
75 the carrier it targeted and not the two beside it.

76 **The signal is not sabotage.** Three measurements point the same way; every matched-protocol
77 number here is comment-stripped on both classes, and the unstripped values, inflated by the comment
78 channel, are in appendix A. First, the classifier *transfers across backdoor generators* at 0.969 pairwise:
79 trained on one attacker’s backdoors it identifies another’s, so what it has learned is not any single
80 model’s sabotage style. Second, it is nearly as strong on **inert** backdoors (0.879) as on working ones
81 (0.935) — if it were detecting sabotage, backdoors that do not function should be much harder, and
82 they are not. Third, the residual carrier is stylistic rather than semantic: `__main__` appears in 0.116
83 of attack samples against **0.000** of honest ones, replicated on the train split (152 versus 1 of 1,209)
84 and untouched by the stripper. None of these is a property of a backdoor. A token-level test [14] does
85 not replicate, and one authorship leak survives the stripper (appendices A, A).

86 **Regenerating one half of a pair does not isolate content.** Under the identical pipeline the human-
87 honest baseline is $B = 0.871$ over $n = 1,444$ pairs, and the routing rule and its bands were fixed
88 before any arm ran: retained signal $\rho = (G - 0.5)/(B - 0.5)$, with $\rho \leq 0.33$ reading as authorship
89 and $\rho > 0.72$ as content. With an open-weight model writing the honest class against the same
90 backdoors ($n = 668$, $B' = 0.874$ on those problems), separation *rises* to $G = 0.962$, $\rho = 1.234$:
91 substituting the honest author replaced one authorship confound with a larger one.

92 **Holding authorship constant needs care about what else changes.** We generate the attack class
93 twice from one honest class. An **edit** arm shown its own honest solution returns near-copies, so only
94 the **independent** arm — shown the problem statement alone, and matching the stored regime of
95 independent programs — changes authorship and nothing else (Table 3, appendix).

96 **Holding authorship constant removes between a third and nine-tenths of the separability, and**
97 **the attack prompt’s framing most of what remains.** With one model writing both classes — one
98 outcome-blind draw per problem, the attack class defined, as in the benchmark, by intent plus passing
99 the public tests (2.2% of these draws fire, the same-generator analogue of the split’s own 80.7%
100 inert class), and B' inert-inclusive over the same problems — $S = 0.751$ (95% CI 0.7295–0.7713;
101 $n = 623$ problem-disjoint pairs) against $B' = 0.8658$, so $\rho = 0.687$ (0.625–0.755, joint bootstrap
102 over problems); a second generator writing both classes under the same rule ($n = 611$) gives
103 $S = 0.5340$ against $B' = 0.8584$, $\rho = 0.095$ (0.044–0.150) on an attack class 99.67% inert against
104 the first generator’s 97.75%. The first interval excludes the collapse band and spans the content
105 line — the pre-registered rule returns *inconclusive between partial and content* — and the second
106 lies inside the collapse band. What survives is lexical: length-only falls to 0.5851 while TF-IDF
107 holds. It is not the trigger: masking numerals and comparison operators leaves S unchanged, while
108 normalising identifiers takes it to 0.700 ($\rho = 0.54$), and the top attack-side features are `def main`
109 and `data[` (67% vs 38.5% of pair sides) — the layout the attack prompt elicits. Holding the prompt
110 constant as well removes most of the rest: with the honest class regenerated under the attack prompt’s
111 framing minus the flaw sentence, $S = 0.6423$ (0.6179–0.6621; $n = 586$), $\rho = 0.394$ (0.326–0.454),
112 and the two honest classes separate from each other at 0.6006. On the second generator neither
113 moves: S with the prompt matched too is 0.5399 ($\rho = 0.113$) and its prompt-only floor is 0.4918, an
114 interval containing chance. The pairs are independent programs (Table 3), and at $n = 623$ B under
115 the identical pipeline is 0.8521 ± 0.0055 , so the drop is not sample size. The strict row — attacks
116 that pass *and* fire, kept by the benchmark’s own retry-and-keep procedure, and the row that follows
117 our own Recommendation 3 — agrees, at low power: $\rho = 0.921$ (0.662–1.089 on the pooled-score
118 0.86) from $S = 0.7652$ (0.7016–0.8000; $n = 99$) against its own $B' = 0.7877$, read through ρ alone
119 because B at that n is 0.7807 ± 0.0248 ; retention is not the artifact, and the edit arm bounds S under
120 near-copies (appendix A). A monitor is inflated the same way. The same zero-shot trusted-monitor
121 prompt scores 0.9579 on the *stored* pairs over the same problems and 0.566 (0.5507–0.5833) on the

headline pairs, a monitor- ρ of 0.145 (0.111–0.182) against the classifier’s 0.687; on the strict pairs it is 0.9437 and 0.818 (0.7630–0.8686), monitor- ρ 0.717 (0.601–0.832), and a monitor that wrote neither class returns the same two ratios (0.151, 0.749). Classifier and monitor are independent pair by pair ($\phi = 0.04$). Our single-shot attacks fire given a pass at 2.5% against the stored attacker’s 67.2%: the stored attack class is a survivor set.

The filter installs a second artifact. Labelling by `backdoor_works` defines the negative class by *attack failure* at the problem level rather than by honesty: 80.7% of the test split’s labelled backdoors do not work (32.8% of those that pass the public tests), and a problem whose backdoor fails contributes its human honest solution to the negative class instead, so the label is fixed before any code is read. Attack problems carry roughly half as many test cases as honest ones (20.95 versus 38.85), consistent with a problem with fewer tests being easier to backdoor undetectably, so difficulty and test count are correlated with the label with no code read at all — which is where the content-blind floor of 0.772 comes from.

The negative control comes free. In five same-generator agentic settings the same three models write both classes, so no authorship confound is possible and any surviving floor is structural (Figure 2, appendix): floors run from 0.504 (bash, on its null) to 0.802 (agentdojo) from character count alone. A per-benchmark floor separates corpora where the label is trivially recoverable from those where it is not.

4 Recommendations

Following Kapoor and Narayanan [10], who treat leakage as a checklist rather than a verdict:

1. **Generate both classes with the same model and the same prompt framing.** If the attack class is model-written, the honest class must be too; as our arm-1 null shows, regenerating only one side substitutes one authorship gap for another.
2. **Report a partial-input floor and a bag-of-words floor beside every monitor number,** each against a permutation null.
3. **Do not define the negative class by attack failure.** Drop the problems whose attacks failed, or sample the negative class independently of attack success.
4. **State the split, the metric and the preprocessing** whenever monitors are compared, and **release the split and the scripts.** Ours accompany this paper under MIT: the split, both selection rules, the retry pool, the scripts, a card and a manifest (appendix A).

5 Limitations

A high floor is evidence about a corpus, not a verdict on results computed on it. Feng et al. [2] warn that the *failure* of a partial-input baseline does not show a dataset is artifact-free; the converse binds us: our floors show the label is recoverable without reading a backdoor [6, 16], not that any monitor exploits that.

The same-generator arm is a later model behind a paid API. The stored attacker, `claude-3-7-sonnet-thinking`, is retired (404 against a live key), so both classes come from a later model of the same lineage, called through a metered commercial API that does not expose its sampling parameters: the arm is not bit-reproducible, and thinking was off although the stored attacker was a thinking model. We do not length-match; the untrusted model chooses its own trajectory length. The stored attacker and the same-generator attacker also differ in backdoor style, not only in authorship, so every ρ here — classifier or monitor — bounds how much of the separability is content and does not partition it; and the two generators are compared on the passes eligibility alone, Haiku 4.5’s strict row being refused at 35 usable pairs against our 50-pair floor.

Scope, provenance and the harness. Whether stored-pair corpora leak this way in general is not established here: five of our seven corpora are one user’s repackaging, and claims are scoped to those releases. Two interpreter-era artifacts fall on opposite classes; the same-generator arm is immune to them, B is not (appendix A).

References

- [1] Basak Demirok, Mucahid Kutlu, and Selin Mergen. MultiAIGCD: A comprehensive dataset for AI generated code detection covering multiple languages, models, prompts, and scenarios. 2025. arXiv:2507.21693v1.
- [2] Shi Feng, Eric Wallace, and Jordan Boyd-Graber. Misleading failures of partial-input baselines. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2019.
- [3] Catherine Ge-Wang, Tyler Crosse, Benjamin Hadad, Joachim Schaeffer, Ram Potham, and Tyler Tracy. Attack selection in agentic AI control evaluations meaningfully decreases safety. 2026. arXiv:2606.06529.
- [4] Ryan Greenblatt, Buck Shlegeris, Kshitij Sachan, and Fabien Roger. AI control: Improving safety despite intentional subversion. 2023. arXiv:2312.06942.
- [5] Hanxi Guo, Siyuan Cheng, Kaiyuan Zhang, Guangyu Shen, and Xiangyu Zhang. CodeMirage: A multi-lingual benchmark for detecting AI-generated and paraphrased source code from production-level LLMs. 2025. arXiv:2506.11059v1. No venue stated.
- [6] Suchin Gururangan, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel R. Bowman, and Noah A. Smith. Annotation artifacts in natural language inference data. In *Proceedings of NAACL-HLT*, 2018.
- [7] Dan Hendrycks, Steven Basart, Saurav Kadavath, Mantas Mazeika, Akul Arora, Ethan Guo, Collin Burns, Samir Puranik, Horace He, Dawn Song, and Jacob Steinhardt. Measuring coding challenge competence with APPS. In *NeurIPS Datasets and Benchmarks Track*, 2021. arXiv:2105.09938. The upstream source of the human-written honest class, distributed as the HuggingFace dataset codeparrot/apps.
- [8] Rogan Inglis. apps-control-arena. HuggingFace dataset `RoganInglis/apps-control-arena`. MIT licence. 5,000 rows; honest class from codeparrot/apps., 2025.
- [9] Anik Jha. Decorrelation is not complementarity: Skill, not lineage, governs trusted-monitor ensembles. 2026. arXiv:2608.16190. Code: <https://github.com/anik-jha/challenger-panels> (MIT).
- [10] Sayash Kapoor and Arvind Narayanan. Leakage and the reproducibility crisis in ML-based science. 2022. arXiv:2207.07048.
- [11] Eugene Koran. apps-control-arena-high-quality. HuggingFace dataset `eugenekoran/apps-control-arena-high-quality`. **No licence declared.** 1,209 train / 1,375 test., 2026.
- [12] Eugene Koran, Yejun Yun, Samantha Tetef, Benjamin Arnav, and Pablo Bernabeu-Pérez. Ensemble monitoring for AI control: Diverse signals outweigh more compute. 2026. arXiv:2605.15377v2.
- [13] Bryce Little. Length penalties make chain-of-thought less monitorable. 2026. arXiv:2607.09786v3.
- [14] Daniil Orel, Dilshod Azizov, Indraneil Paul, Yuxia Wang, Iryna Gurevych, and Preslav Nakov. AICD bench: A challenging benchmark for AI-generated code detection. 2026. arXiv:2602.02079v1. No venue stated on the abs page.
- [15] Lucas Pinto. Calibration-family overfit: Why trusted sabotage monitors don’t transfer across lineages. In *COLM 2026 Workshop on AI Measurement Science*, 2026. arXiv:2607.06596v3.
- [16] Adam Poliak, Jason Naradowsky, Aparajita Haldar, Rachel Rudinger, and Benjamin Van Durme. Hypothesis only baselines in natural language inference. In *Proceedings of the Seventh Joint Conference on Lexical and Computational Semantics (*SEM)*, 2018.

- [17] Aditya Sinha. control_arena_{bash,shade,iac,rogue_eval,agentdojo}. HuggingFace datasets adityaasinha28/control_arena_*. **No licence declared.** 1,500 rows each., 2026.
- [18] Ashwin Sreevatsa, Sebastian Prasanna, and Cody Rushing. Basic legibility protocols improve trusted monitoring. 2026. arXiv:2602.10153.
- [19] Mikhail Terekhov, Zhen Ning David Liu, Caglar Gulcehre, and Samuel Albanie. Control tax: The price of keeping AI in check. 2025. arXiv:2506.05296v3.
- [20] Mikhail Terekhov, Alexander Panfilov, Daniil Dzenhaliou, Caglar Gulcehre, Maksym Andriushchenko, Ameya Prabhu, and Jonas Geiping. Adaptive attacks on trusted monitors subvert AI control protocols. In *International Conference on Learning Representations (ICLR)*, 2026. arXiv:2510.09462v2.

A Supplementary tables and figures

The schema. Table 1 is the record schema of RoganInglis/apps-control-arena [8]. The attack class carries model and source, naming the system that produced each backdoor; the honest class carries no such field, because there is no generator to name. The two APPS-backdoor corpora we use are two views of this one artifact: every problem in the filtered release [11] is in [8], byte-identical except a re-validated backdoor_works. Among the agentic corpora we exclude AgentHarm, whose author is an invited speaker at this venue.

Table 1: The confound is legible in the dataset schema, with no measurement at all. The attack class records which model wrote it; the honest class has no generator field. Verified against the shard schema of RoganInglis/apps-control-arena.

struct	sub-fields
solutions	code, compiles, passes_tests
backdoors	backdoor_code, backdoor_input, backdoor_works, compiles, model, passes_tests, source

The protocol $\times$ baseline table. The no-code floor is cross-validated *within* the test split because the released train/val split is paired and therefore metadata-balanced; its permutation null shuffles labels for fitting and scoring alike. Table 2 is the full measurement behind §3. Rows are not comparable across protocols; the preprocessing column says whether comments were stripped before the baseline read the code. Pairwise accuracies are given to three decimals because their resolution over 2,641 pairs is one pair. Intervals in §3 are bootstraps over problems (2,000 draws) and nulls are 200 label permutations; the held-out no-code row is barely above its own null (0.716 against a 95th percentile of 0.7023, $p = 0.030$), so the no-code claim rests on the cross-validated protocol.

Comment stripping on the matched-protocol numbers. The matched-pair and transfer measurements were first made with comments in. Every value in the authorship paragraph of §3 was re-measured with the same stripper applied to both classes, over five problem-disjoint splits (each transfer row is one fixed split). With comments kept the same quantities were transfer 0.994, working 0.979, inert 0.955, all pairs 0.962 and surface-feature AUROC 0.8306; stripped they are 0.969, 0.935, 0.879, 0.897 and 0.7259. Every one was inflated by the comment channel, which is this paper’s thesis applied to itself. The #-presence marker, 0.685 against 0.176 before stripping, closes to 0.007 against 0.008 after and is dropped; __main__ is unchanged. One authorship leak survives the stripper because string literals are not comments: in the comment-stripped train split an author byline (Author :) survives in 23 of 1,209 honest samples against 2 attack samples, and a triple-quoted block in 43 against 9 — too rare to drive 0.806, but a literal record of authorship. Stripped, the matched-protocol TF-IDF AUROC is 0.8560, beside $B = 0.871$ and the Koran-split 0.806, all three on stripped code. The other cross-generator transfer rows, stripped, are 0.866 (o3-mini to claude-3-7-sonnet, 82 test pairs), 0.844 (the reverse) and 0.764 (o1-mini to gemini-2.0-flash).

Table 2: What the field’s preprocessing leaves behind. Every row is measured under the protocol named in column one; rows are not comparable across protocols. *Prep.* says whether comments were stripped before the baseline read the code; *no code* means the baseline never reads the program. Pairwise accuracies are given to three decimals, their resolution over the pair count.

Protocol	Baseline	Value	Prep.	Against
<i>AUROC</i>				
Koran’s exact split	TF-IDF word 1–2gram	0.8063	stripped	above 20 of 24
same	TF-IDF word 1–2gram	0.8914	kept	above 22 of 24
same, reads no code	difficulty + #tests + $ q $	0.7715 ± 0.0389	no code	above 20 of 24
same, held-out slice	difficulty + #tests + $ q $	0.7163	no code	above 18 of 24
same	code length only	0.4726	stripped	chance
same	#-count detector	0.4983	stripped	chance
same	#-count detector	0.6957	kept	—
Raw artifact, problem level	all metadata	0.8917	no code	—
same	difficulty only	0.6823	no code	—
same	<code>solution_passes_tests</code>	0.7881	no code	—
Matched pairs (P3)	all metadata	0.5000	no code	exactly chance by construction
<i>Pairwise accuracy</i>				
DSPy matched, problem-disjoint	TF-IDF word 1–2gram	0.897	stripped	tutorial monitor 0.69–0.81
working backdoors only	TF-IDF word 1–2gram	0.935	stripped	—
inert backdoors only	TF-IDF word 1–2gram	0.879	stripped	the signal is not sabotage
transfer, Control Tax → GPT-4-era set	TF-IDF word 1–2gram	0.969	stripped	—
all pairs, comments kept	TF-IDF word 1–2gram	0.962	kept	inflated by the comment channel
working, comments kept	TF-IDF word 1–2gram	0.979	kept	inflated by the comment channel
inert, comments kept	TF-IDF word 1–2gram	0.955	kept	inflated by the comment channel
transfer, comments kept	TF-IDF word 1–2gram	0.994	kept	inflated by the comment channel
same, raw artifact	code length only	0.826	kept	—

256 **Pair similarity, and why the same-generator control needs two arms.** An attack class built
257 by editing its own honest solution changes pair similarity as well as authorship (Table 3); the
258 independent arm sits in the stored benchmark’s regime and if anything slightly further from it, so it
259 cannot manufacture a low S by making the classes alike.

Table 3: Character similarity of honest/attack pairs, comments stripped. The edit arm’s classes are near-copies; the independent arm reproduces the stored benchmark’s regime, where the two classes are largely independent programs.

Attack class	mean sim.	> 0.95	< 0.60	changes vs. the benchmark
stored benchmark	0.606	3.4%	45.4%	— (the regime to match)
edit (shown its own solution)	0.928	64.2%	4.0%	authorship <i>and</i> similarity
independent (problem only)	0.659	3.4%	34.8%	authorship only

260 **Robustness of the same-generator rows.** Intervals in the Gate S paragraph of §3 are bootstraps
261 over problems, resampling S and B' jointly for ρ ; this replaced a ± 2 fold-sd rule before the retry
262 pool was read. Every ρ point we quote is the ratio of the fold-mean AUROCs and every joint interval
263 beside it is the bootstrap of the pooled out-of-fold ρ , which is the same quantity computed two ways:
264 the pooled points are 0.390 (prompt-matched headline), 0.662 (prompt-matched strict) and 0.097
265 (the second generator). Five fold seeds move the strict row’s ρ over 0.85–0.92. Retention is not the
266 artifact: survivors and non-survivors separate at only 0.62, and S with non-firing attempts on the
267 same problems is 0.7595 against 0.7687. The edit arm bounds S under near-copies (similarity 0.928,
268 Table 3): 0.6266 over 427 pairs, $\rho = 0.371$. In arm 1, where only the honest half was regenerated,
269 TF-IDF gains +0.088 while the length carrier collapses to 0.489. Half the gap between our single-shot

270 firing rate and the stored attacker’s is our own attack prompt: 52.0% of our test-passing draws offer a
 271 public test input as the trigger, which cannot diverge from the honest solution, and those fire 0.15%
 272 against 5.13% for novel triggers; no stored working backdoor has such a trigger (0 of 3,188). The
 273 rest is retry-and-keep.

274 **The prompt-matched arm under the masks, and the second generator.** Under prompt matching
 275 the identifier mask — the only mask of §3 that moved S — takes the headline row to 0.6264
 276 ($\rho = 0.342$) where it takes the original honest class on the same 586 problems to 0.7046: most
 277 of what identifier normalisation removed is what prompt matching removes. That masked row is
 278 the one that decides whether the residue is structure or naming, and under prompt matching it is
 279 band-ambiguous: its joint interval, 0.280–0.394, puts 40% of its mass in the collapse band. A
 280 character model (3–5-grams, same classifier and folds) retains more than the word tokeniser, 0.6817
 281 ($\rho = 0.493$), so the drop is not an artefact of tokenisation. On the strict pairs the matched row is
 282 $S = 0.7113$ ($n = 96$, $\rho = 0.697$, 0.494–0.856). The second generator’s pairs are more self-similar
 283 than the first’s (0.724 mean against 0.659), so S was recomputed inside similarity strata: it stays
 284 0.4737–0.5142 in every stratum, including the 0.451-similarity stratum, which is further from its
 285 honest class than the stored benchmark’s own pairs, while B' holds 0.7884–0.8334. The collapse is
 286 not bought with near-copies. Its prompt-matched rows move S to 0.5399 ($\rho = 0.113$, $n = 562$) and
 287 put the prompt-only floor at 0.4918, an interval containing chance. The mechanism is visible in the
 288 programs: the layout vocabulary the first generator’s residue is made of — `def main, sys.stdin,`
 289 `data[` — is present in 0.000 of the second generator’s programs in *all three* of its classes, which read
 290 standard input with `input(` in 1.000 of programs on every side, so there is no layout difference left
 291 for the classifier to read.

292 **The monitor rows, and what bounds the second monitor.** The monitor that wrote neither class
 293 scores 0.5577 on the same-generator headline pairs and 0.8824 on the stored pairs over the same
 294 problems, 0.7871 and 0.8832 on the strict pairs. Its headline rows are a bound rather than a point:
 295 a 512-token reply ceiling left 628 of 1,246 same-generator replies unparseable, and on the stored
 296 side the missing share differs by class (50.1% honest against 23.4% attack), a bias of unknown sign.
 297 Re-running the strict pairs with a 1,024-token ceiling prices it: missingness falls to 2 of 198 and
 298 monitor- ρ moves from 0.749 to 0.766, inside the bootstrap interval.

299 **The two generators’ funnel.** Both generators were run through the identical pipeline. Our honest
 300 draws pass the public tests at 0.7701 and 0.7278 of 1,444 problems; the second generator’s attack
 301 draws pass at 689 of 1,444 and 35 of those fire, below the 50-pair floor `gate_s_samegen.py` requires,
 302 which is why it has no strict row and the two generators are compared on the `passes` eligibility only.
 303 The two compared rows are not equally inert either: 2 of the 611 attacks in the second generator’s row
 304 fire (99.67% inert) against 14 of 623 (97.75%), the other 33 of its firing draws sitting on problems
 305 where its own honest program failed the tests and which Gate S therefore drops. The single-shot
 306 firing rate is a property of the protocol rather than of either model: 52.0% of test-passing draws offer
 307 a public test input as their trigger.

308 **The five same-generator agentic settings.** Figure 2 shows each setting’s content-blind floor against
 309 its own permutation null.

310 **The release.** The supplementary bundle holds the same-generator split (1,444 honest programs,
 311 1,112 edit-arm and 1,112 independent-arm attack programs written by one model, plus every one
 312 of the 11,565 draws of the independent arm’s retry pool), both eligibility rules as pair lists under an
 313 outcome-blind selection (623 passing and 14 firing independent-arm pairs) and under the benchmark’s
 314 own survivor selection (1,046 and 99, or 106 counting a crash on the trigger as firing), the 24 scripts
 315 behind the measurements — generation for both generators, the test harness, the trigger verifier, the
 316 selection rules, the prompt-matched arm, the monitor rows and each remaining measurement script —
 317 a dataset card, and a manifest of row counts and checksums. The licence, MIT, is a file inside the
 318 bundle and covers the data and the scripts alike.

319 **The test harness and the pass flag.** Gate S requires executing generated code, so we built a runner
 320 that feeds each shipped input case to a solution under per-case timeouts and address-space limits.
 321 It agrees with the artifact’s own `solution_passes_tests` on 95.84% of the analysis pool and is
 322 strictly conservative there — it never credits a solution the artifact rejects (0 such cases). We use




Figure 2: **The structural floor in five same-generator agentic settings.** Each setting’s content-blind classifier — six count features, no content read — scored *within task* against its own permutation null. No authorship confound is possible here: claude-sonnet-4, gemini-2.5-pro and gpt-5 write both classes in all five. *bash* sits on its null at 0.504; *agentdojo* reaches 0.802 from character count alone. **The task-group count beside each setting matters:** *shade* and *iac* average over five task groups each while *bash* averages over 247, so these floors are not equally well determined. We report the within-task column because our own pooled numbers partly broke on task composition, and we broke them ourselves. The permutation nulls shown were computed with a pooled estimator rather than the folded within-task statistic; they bound the chance level but are not a like-for-like null for each dot. These corpora declare no licence; we redistribute nothing and report aggregates only.

323 *our own* flag for both classes: a checker stricter on one class than the other would manufacture
 324 exactly the artifact this paper documents. The headline does not depend on that choice — under the
 325 artifact’s flag $B = 0.871$ at $n = 1,444$, under ours $B = 0.873$ at $n = 1,406$, within one standard
 326 deviation of each other. Our release carries one MIT-licensed file of published monitor scores, with its
 327 notice; the monitor scores in Figure 1 are runs/x0/monitor-summary.json of that release. Two
 328 interpreter-era artifacts in the upstream solutions fall on opposite classes: a `stdin` without `.buffer`,
 329 fixed, and `fractions.gcd`, removed in Python 3.9, not fixable.

330 **The token test, run and failed.** We attempted to confirm the token-level mechanism of [14] in our
 331 setting and it does not replicate: separation -0.013 — the wrong sign — at permutation $p = 0.956$.
 332 The reason is register. Their human class is general-purpose repository code, whereas *both* of our
 333 classes are competitive-programming code, so an identifier axis that separates general human code
 334 from model code does not separate these two populations. Only 8 of their 10 canonical tokens are in
 335 our vocabulary, so the test is underpowered; but the point estimate has the wrong sign, which is not
 336 what an underpowered true effect looks like. We report it because a reader who knows that work may
 337 run exactly this test.

338 **A third channel, noted rather than developed.** In the same five agentic settings a fifteen-word
 339 regular expression over the untrusted policy’s own scratchpad reaches 0.76–0.82, because prompted
 340 attack policies restate their side task in the reasoning the monitor is given. We note this only to bound
 341 the claim: the instrument belongs to [13], which validates it properly against an LLM judge and a
 342 length-matched null, and we defer to that treatment.