
What Survives the Filter? Trivial Baselines on Stored Honest/Attack Pairs in AI-Control Monitoring

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Trusted monitoring in AI control is evaluated on stored honest/attack pairs. The
2 field already knows these pairs leak: it filters to working solutions and strips
3 comments. We ask what that preprocessing leaves behind. On the field’s own
4 filtered, comment-stripped split, a bag-of-words classifier scores above roughly
5 four-fifths of the open-weight monitors ranked on it, and a classifier that reads no
6 code at all — only problem metadata — scores above most of them. We trace the
7 residue to authorship rather than sabotage: the honest class is human-written while
8 every attack is model-written. The signal transfers across backdoor generators, is
9 nearly as strong on inert backdoors as on working ones, and survives the field’s
10 comment stripper. Five same-generator agentic settings, where no authorship
11 confound is possible, carry a second, structural floor. Regenerating only the
12 honest half does not isolate content: with an open-weight model writing honest
13 solutions against the stored backdoors, separability *rises*. Holding authorship
14 constant requires one model to write both classes; we report that design and its
15 single-shot yield, which shows how heavily the stored attack class has been filtered.
16 We recommend that stored-pair evaluations report partial-input and bag-of-words
17 floors against permutation nulls, use same-generator honest samples, and drop inert
18 backdoors from both classes.

19 1 Introduction

20 Trusted monitoring is the backbone protocol of AI control: a weaker trusted model reads an untrusted
21 model’s output and flags sabotage [4]. It is developed, tuned and ranked almost entirely on *stored*
22 honest/attack pairs — fixed corpora of solutions, some carrying a planted backdoor [12, 19, 9]. The
23 field already knows these pairs leak. Koran et al. keep only working honest solutions and working
24 backdoors and strip every comment, with the rationale stated [12]; Control Tax reports the pass-tests
25 pre-filter and appendices “Monitoring failed backdoors” [19]; [18] is an entire paper on the comment
26 channel. The sharpest treatment is Jha’s, whose released harness ships a `-leak-check` flag with two
27 trivial baselines already inside it and concludes that stripping is mandatory rather than cosmetic: a
28 #-count detector falls from 0.6957 to 0.4983 once comments go, and code length alone sits at 0.4726
29 [9]. This paper starts from that work, not against it.

30 What nobody has measured is what the preprocessing leaves behind. On the field’s own filtered,
31 comment-stripped split, a bag-of-words classifier reaches AUROC 0.806 — above 20 of the 24
32 open-weight monitors ranked on that same split, whose median is 0.609 — and a classifier that reads
33 no code at all, only problem difficulty, test count and question length, reaches 0.772, above 20.

34 The residue is authorship, not sabotage. Honest solutions come from APPS [7] and are human-
35 written; every backdoor is model-written. The classifier transfers across backdoor generators and is
36 nearly as strong on **inert** backdoors as on working ones — neither is a property of a backdoor. That

bag-of-words separates human-written from model-written code is established out-of-field [14]; the known-ness of the mechanism is the argument, not a dent in it. The neighbouring field already applies the control missing here — CodeMirage gates its two classes on line count and character length [5], MultiAIGCD pairs to shared problem definitions [1] — so unmatched pairs are a defect by a published standard, not one we invented. Adjacent control work attacks other axes: [20] subverts monitors with injected text, [3] shows attacker-side selection inflates safety estimates, and [15] removes transcript length as a nuisance covariate, where we publish a content-blind AUROC as a per-benchmark floor.

We contribute (i) a protocol $\times$ baseline measurement of what survives the field’s preprocessing; (ii) evidence that the survivor is authorship, including a negative result: regenerating one half of a stored pair swaps the confound rather than removing it; (iii) a second selection artifact the filter installs by defining the negative class through attack failure; and (iv) a preprocessing checklist and a released split. **Scope:** this concerns monitor training, selection and benchmarking on *stored* pairs. ControlArena’s live evaluations generate both classes and are not implicated.

2 Setup

The confound is legible in the schema. The canonical stored-pair artifact for APPS-based control evaluations [8] records `model` and `source` on every backdoor and no generator field on any honest solution (Table 1, appendix): the honest class is taken from APPS [7] and names no model; every backdoor is model-written.

Artifacts and provenance. We use seven public corpora; their per-class provenance is the object of study: in the APPS family the honest class is human and the attack class is one of ten named models, while in the five agentic settings the same three models write *both* classes. **Six of the seven, from two publishers, declare no licence** [11, 17]; we report only aggregates over them and redistribute no corpus.

Evaluation surface, and naming protocols. Our headline protocol is the filtered, comment-stripped split of [12]: 1,209 train and 1,375 test items (172 attack, 1,203 honest), labelled by `backdoor_works`, reconstructed with the released code of [9]. It is the surface on which 24 open-weight monitors have published per-item scores. The raw artifact, this split, pooled evaluation, matched pairs and within-task evaluation are five populations, and pairwise accuracy and AUROC two metrics; every comparison below names its split, metric and preprocessing, and no value is placed beside a count from another protocol. Generated code is scored by our own runner, which agrees with the artifact’s `solution_passes_tests` on 95.84% of the analysis pool and is applied to both classes (appendix A).

3 Results

Table 2 (appendix) reports what a trivial classifier recovers under each protocol. The fix works: code length is driven to chance on the filtered split (0.4726) and the #-count detector falls from 0.6957 to 0.4983. What it does not reach, on that same split, is a bag-of-words floor of 0.806 (95% CI 0.7704–0.8409 by bootstrap over problems; permutation null 95th percentile 0.5465, $p < 0.005$), above 20 of the 24 monitors by point estimate and 18 of them with a paired-bootstrap interval excluding zero, and a no-code floor of 0.772 (CI 0.7319–0.8005; null 0.5545, $p < 0.005$), above 20 and 17 respectively (Figure 1): the preprocessing removed the carrier it targeted and not the two beside it.

The signal is authorship, not sabotage. Three measurements point the same way; every matched-protocol number here is comment-stripped on both classes, and the unstripped values, inflated by the comment channel, are in appendix A. First, the classifier *transfers across backdoor generators* at 0.969 pairwise: trained on one attacker’s backdoors it identifies another’s, so what it has learned is not any single model’s sabotage style. Second, it is nearly as strong on **inert** backdoors (0.879) as on working ones (0.935) — if it were detecting sabotage, backdoors that do not function should be much harder, and they are not. Third, the residual carrier is stylistic rather than semantic: `__main__` appears in 0.116 of attack samples against **0.000** of honest ones, replicated on the train split (152 versus 1 of 1,209) and untouched by the stripper. None of these is a property of a backdoor. All

of them are properties of who wrote the file. A fourth, token-level test [14] does not replicate here (separation -0.013 , permutation $p = 0.956$; appendix A), a register effect. One authorship leak even survives the field’s own comment stripper, because string literals are not comments: in the comment-stripped train split an author byline (Author :) survives in 23 of 1,209 honest samples against 2 attack samples, and a triple-quoted block in 43 against 9 — too rare to drive 0.806, but a literal record of authorship.

Regenerating one half of a pair does not isolate content. Under the identical pipeline the human-honest baseline is $B = 0.871$ over $n = 1,444$ pairs, and the routing rule and its bands were fixed before any arm ran: retained signal $\rho = (G - 0.5)/(B - 0.5)$, with $\rho \leq 0.33$ reading as authorship and $\rho > 0.72$ as content. With an open-weight model writing the honest class against the same backdoors ($n = 668$, $B' = 0.874$ on those problems), separation *rises* to $G = 0.962$, $\rho = 1.234$, some 4.78 standard deviations above B' ; TF-IDF gains $+0.088$ while the length carrier collapses to 0.489. The new gap is lexical, not structural: substituting the honest author replaced one authorship confound with a larger one, and any control that regenerates one side of a stored pair inherits the same defect.

Holding authorship constant needs care about what else changes. We generate the attack class twice from one honest class: an **edit** arm shown its own honest solution and an **independent** arm shown only the problem statement. The stored benchmark’s pairs are largely independent programs (mean character similarity 0.469); the edit arm’s are near-copies (0.924) while the independent arm matches the stored regime (0.415; Table 3, appendix). Only the independent arm changes authorship alone — a low S from the edit arm could mean only that we made the two classes the same text — so we route on it and report the edit arm only as a bound.

The same-generator control did not land, and the authorship reading is inference. Holding authorship constant requires one model to write both classes; our attack generations converted at roughly 2% into pairs that both pass the public tests and fire on their trigger — 40 of 1,112 in the edit arm and 10 of 600 in the independent arm, against the 50-pair floor our analysis script requires — so we report no S . The authorship reading is inference from the transfer, inert-backdoor and style evidence together with the five same-generator settings. The yield gap is itself a measurement: ControlArena attack policies resample and keep survivors, so the stored attack class is heavily filtered relative to any single-shot rate. One rule was changed before the retry batch was read, and we say so: the ± 2 fold-sd interval first coded for S is far wider than a 95% interval over problems and biased toward “inconclusive”, so S now routes on a bootstrap-over-problems interval.

The filter installs a second artifact. Labelling by `backdoor_works` defines the negative class by *attack failure* rather than by honesty: 80.7% of the test split’s labelled backdoors do not work (32.8% of those that pass the public tests) and are relabelled honest, so most of the “honest” class consists of failed attacks. Attack problems carry roughly half as many test cases as honest ones (20.95 versus 38.85), consistent with a problem with fewer tests being easier to backdoor undetectably, so difficulty and test count are correlated with the label before any code is read — which is where the content-blind floor of 0.772 comes from.

The negative control comes free. In five same-generator agentic settings the same three models write both classes, so no authorship confound is possible and any surviving floor is structural (Figure 2, appendix). `bash` sits on its permutation null at 0.504 while `agentdojo` reaches 0.802 from character count alone. A per-benchmark floor therefore separates corpora where the label is trivially recoverable from those where it is not. **Two mechanisms:** authorship on the APPS pairs, structure on the agentic settings, and neither is visible without running the baseline.

4 Recommendations

Following Kapoor and Narayanan [10], who treat leakage as a checklist rather than a verdict:

1. **Generate both classes with the same model.** If the attack class is model-written, the honest class must be too; as our arm-1 null shows, regenerating only one side substitutes one authorship gap for another. This is precedent [5, 1], not invention.




Figure 1: **Two floors against the 24 open-weight monitors ranked on the same split.** All values are AUROC on Koran et al.’s 1,375-item test split (172 attack / 1,203 honest), comments stripped. Monitor scores from [9] (MIT); median 0.609, max 0.937. **Two caveats.** (i) The 24 monitors are zero-shot prompted while the bag-of-words classifier is *trained* on the benchmark’s own paired train split — this is a statement about what the corpus makes recoverable, in the sense of [6], not a claim that bag-of-words is a better monitor. (ii) The no-code floor is fitted *within* the test split by 5×5 cross-validation, because the released train/val split is paired and therefore metadata-balanced, so its permutation null shuffles labels for fitting and scoring alike; on a single held-out slice of the same items it is 0.716 (CI 0.6476–0.7796; $p = 0.030$ against its own null), above 18 of 24, 12 with a paired interval excluding zero.

2. **Report a partial-input floor and a bag-of-words floor beside every monitor number,** each against a permutation null.
3. **Do not define the negative class by attack failure.** Drop inert backdoors from both classes rather than relabelling a failed attack as honest.
4. **State the split, the metric and the preprocessing** whenever monitors are compared, and **release the split and the scripts.** Ours accompany this paper as supplementary material under MIT: the same-generator split (1,444 honest, 1,112 edit-arm and 1,112 independent-arm programs from one model, plus all 11,565 retry-pool draws), pair lists under both selection rules — outcome-blind (623 passing, 14 firing independent-arm pairs) and the benchmark’s own survivor selection (1,046 and 99) — a dataset card and a manifest (appendix A).

5 Limitations

A high floor is evidence about a corpus, not a verdict on results computed on it. Feng et al. [2] warn that the *failure* of a partial-input baseline does not show a dataset is artifact-free; the converse binds us: our floors show the label is recoverable without reading a backdoor, in the sense of Gururangan et al. [6] and Poliak et al. [16], not that any monitor exploits that (Figure 1, caveat i).

The same-generator arm is a later model behind a paid API. The stored attacker, `claude-3-7-sonnet-thinking`, is retired (404 against a live key), so both same-generator classes come from a later model of the same lineage, called through a metered commercial API. That API does not expose its sampling parameters, so the arm is not bit-reproducible as the open-weight arms are, and thinking was off although the stored attacker was a thinking model. That the remaining cross-model difference is small is an inference from our 0.969 transfer and Orel et al. [14]’s generator-independence result. We do not length-match: in deployment the untrusted model chooses its trajectory length. The authorship reading is therefore inference (§3).

Scope, provenance and the harness. Whether stored-pair corpora leak this way in general is not established here. Five of our seven corpora are one user’s repackaging; claims are scoped to those releases. Two interpreter-era artifacts fall on opposite classes (a patched `stdin` lacking `.buffer` penalised the model class, fixed; `fractions.gcd`, removed in Python 3.9, penalises old human code); the same-generator arm is immune, *B* is not.

References

- [1] Basak Demirok, Mucahid Kutlu, and Selin Mergen. MultiAIGCD: A comprehensive dataset for AI generated code detection covering multiple languages, models, prompts, and scenarios. 2025. arXiv:2507.21693v1.
- [2] Shi Feng, Eric Wallace, and Jordan Boyd-Graber. Misleading failures of partial-input baselines. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2019.
- [3] Catherine Ge-Wang, Tyler Crosse, Benjamin Hadad, Joachim Schaeffer, Ram Potham, and Tyler Tracy. Attack selection in agentic AI control evaluations meaningfully decreases safety. 2026. arXiv:2606.06529.
- [4] Ryan Greenblatt, Buck Shlegeris, Kshitij Sachan, and Fabien Roger. AI control: Improving safety despite intentional subversion. 2023. arXiv:2312.06942.
- [5] Hanxi Guo, Siyuan Cheng, Kaiyuan Zhang, Guangyu Shen, and Xiangyu Zhang. CodeMirage: A multi-lingual benchmark for detecting AI-generated and paraphrased source code from production-level LLMs. 2025. arXiv:2506.11059v1. No venue stated.
- [6] Suchin Gururangan, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel R. Bowman, and Noah A. Smith. Annotation artifacts in natural language inference data. In *Proceedings of NAACL-HLT*, 2018.
- [7] Dan Hendrycks, Steven Basart, Saurav Kadavath, Mantas Mazeika, Akul Arora, Ethan Guo, Collin Burns, Samir Puranik, Horace He, Dawn Song, and Jacob Steinhardt. Measuring coding challenge competence with APPS. In *NeurIPS Datasets and Benchmarks Track*, 2021. arXiv:2105.09938. The upstream source of the human-written honest class, distributed as the HuggingFace dataset codeparrot/apps.
- [8] Rogan Inglis. apps-control-arena. HuggingFace dataset `RoganInglis/apps-control-arena`. MIT licence. 5,000 rows; honest class from codeparrot/apps., 2025.
- [9] Anik Jha. Decorrelation is not complementarity: Skill, not lineage, governs trusted-monitor ensembles. 2026. arXiv:2608.16190. Code: <https://github.com/anik-jha/challenger-panels> (MIT).
- [10] Sayash Kapoor and Arvind Narayanan. Leakage and the reproducibility crisis in ML-based science. 2022. arXiv:2207.07048.
- [11] Eugene Koran. apps-control-arena-high-quality. HuggingFace dataset `eugenekoran/apps-control-arena-high-quality`. **No licence declared.** 1,209 train / 1,375 test., 2026.
- [12] Eugene Koran, Yejun Yun, Samantha Tetef, Benjamin Arnav, and Pablo Bernabeu-Pérez. Ensemble monitoring for AI control: Diverse signals outweigh more compute. 2026. arXiv:2605.15377v2.
- [13] Bryce Little. Length penalties make chain-of-thought less monitorable. 2026. arXiv:2607.09786v3.
- [14] Daniil Orel, Dilshod Azizov, Indraneil Paul, Yuxia Wang, Iryna Gurevych, and Preslav Nakov. AICD bench: A challenging benchmark for AI-generated code detection. 2026. arXiv:2602.02079v1. No venue stated on the abs page.
- [15] Lucas Pinto. Calibration-family overfit: Why trusted sabotage monitors don’t transfer across lineages. In *COLM 2026 Workshop on AI Measurement Science*, 2026. arXiv:2607.06596v3.
- [16] Adam Poliak, Jason Naradowsky, Aparajita Haldar, Rachel Rudinger, and Benjamin Van Durme. Hypothesis only baselines in natural language inference. In *Proceedings of the Seventh Joint Conference on Lexical and Computational Semantics (*SEM)*, 2018.

- [17] Aditya Sinha. control_arena_{bash,shade,iac,rogue_eval,agentdojo}. HuggingFace datasets adityaasinha28/control_arena_*. **No licence declared.** 1,500 rows each., 2026.
- [18] Ashwin Sreevatsa, Sebastian Prasanna, and Cody Rushing. Basic legibility protocols improve trusted monitoring. 2026. arXiv:2602.10153.
- [19] Mikhail Terekhov, Zhen Ning David Liu, Caglar Gulcehre, and Samuel Albanie. Control tax: The price of keeping AI in check. 2025. arXiv:2506.05296v3.
- [20] Mikhail Terekhov, Alexander Panfilov, Daniil Dzenhaliou, Caglar Gulcehre, Maksym Andriushchenko, Ameya Prabhu, and Jonas Geiping. Adaptive attacks on trusted monitors subvert AI control protocols. In *International Conference on Learning Representations (ICLR)*, 2026. arXiv:2510.09462v2.

A Supplementary tables and figures

The schema. Table 1 is the record schema of RoganInglis/apps-control-arena [8]. The attack class carries model and source, naming the system that produced each backdoor; the honest class carries no such field, because there is no generator to name. The two APPS-backdoor corpora we use are two views of this one artifact: every problem in the filtered release [11] is in [8], byte-identical except a re-validated backdoor_works. Among the agentic corpora we exclude AgentHarm, whose author is an invited speaker at this venue.

Table 1: The confound is legible in the dataset schema, with no measurement at all. The attack class records which model wrote it; the honest class has no generator field. Verified against the shard schema of RoganInglis/apps-control-arena.

struct	sub-fields
solutions	code, compiles, passes_tests
backdoors	backdoor_code, backdoor_input, backdoor_works, compiles, model, passes_tests, source

The protocol $\times$ baseline table. Table 2 is the full measurement behind §3. Rows are not comparable across protocols; the preprocessing column says whether comments were stripped before the baseline read the code. Pairwise accuracies are given to three decimals because their resolution over 2,641 pairs is one pair. Intervals in §3 are bootstraps over problems (2,000 draws) and nulls are 200 label permutations; the held-out no-code row is barely above its own null (0.716 against a 95th percentile of 0.7023, $p = 0.030$), so the no-code claim rests on the cross-validated protocol.

Comment stripping on the matched-protocol numbers. The matched-pair and transfer measurements were first made with comments in. Every value in the authorship paragraph of §3 was re-measured with the same stripper applied to both classes, over five problem-disjoint splits (each transfer row is one fixed split). With comments kept the same quantities were transfer 0.994, working 0.979, inert 0.955, all pairs 0.962 and surface-feature AUROC 0.8306; stripped they are 0.969, 0.935, 0.879, 0.897 and 0.7259. Every one was inflated by the comment channel, which is this paper’s thesis applied to itself. The #-presence marker, 0.685 against 0.176 before stripping, closes to 0.007 against 0.008 after and is dropped; `__main__` is unchanged. Stripped, the matched-protocol TF-IDF AUROC is 0.8560, beside $B = 0.871$ and the Koran-split 0.806, all three on stripped code. The other cross-generator transfer rows, stripped, are 0.866 (o3-mini to claude-3-7-sonnet, 82 test pairs), 0.844 (the reverse) and 0.764 (o1-mini to gemini-2.0-flash).

Pair similarity, and why the same-generator control needs two arms. An attack class built by editing its own honest solution changes pair similarity as well as authorship (Table 3); the independent arm sits in the stored benchmark’s regime and if anything slightly further from it, so it cannot manufacture a low S by making the classes alike.

The five same-generator agentic settings. Figure 2 shows each setting’s content-blind floor against its own permutation null.

Table 2: What the field’s preprocessing leaves behind. Every row is measured under the protocol named in column one; rows are not comparable across protocols. *Prep.* says whether comments were stripped before the baseline read the code; *no code* means the baseline never reads the program. Pairwise accuracies are given to three decimals, their resolution over the pair count.

Protocol	Baseline	Value	Prep.	Against
<i>AUROC</i>				
Koran’s exact split	TF-IDF word 1–2gram	0.8063	stripped	above 20 of 24
same	TF-IDF word 1–2gram	0.8914	kept	above 22 of 24
same, reads no code	difficulty + #tests + $ q $	0.7715 ± 0.0389	no code	above 20 of 24
same, held-out slice	difficulty + #tests + $ q $	0.7163	no code	above 18 of 24
same	code length only	0.4726	stripped	chance
same	#-count detector	0.4983	stripped	chance
same	#-count detector	0.6957	kept	—
Raw artifact, problem level	all metadata	0.8917	no code	—
same	difficulty only	0.6823	no code	—
same	<code>solution_passes_tests</code>	0.7881	no code	—
Matched pairs (P3)	all metadata	0.5000	no code	exactly chance by construction
<i>Pairwise accuracy</i>				
DSPy matched, problem-disjoint	TF-IDF word 1–2gram	0.897	stripped	tutorial monitor 0.69–0.81
working backdoors only	TF-IDF word 1–2gram	0.935	stripped	—
inert backdoors only	TF-IDF word 1–2gram	0.879	stripped	the signal is not sabotage
transfer, Control Tax → GPT-4-era set	TF-IDF word 1–2gram	0.969	stripped	—
all pairs, comments kept	TF-IDF word 1–2gram	0.962	kept	inflated by the comment channel
working, comments kept	TF-IDF word 1–2gram	0.979	kept	inflated by the comment channel
inert, comments kept	TF-IDF word 1–2gram	0.955	kept	inflated by the comment channel
transfer, comments kept	TF-IDF word 1–2gram	0.994	kept	inflated by the comment channel
same, raw artifact	code length only	0.826	kept	—

Table 3: Character similarity of honest/attack pairs, comments stripped. The edit arm’s classes are near-copies; the independent arm reproduces the stored benchmark’s regime, where the two classes are largely independent programs.

Attack class	mean sim.	> 0.95	< 0.60	changes vs. the benchmark
stored benchmark	0.469	2.5%	60.8%	— (the regime to match)
edit (shown its own solution)	0.924	68.7%	4.3%	authorship <i>and</i> similarity
independent (problem only)	0.415	1.0%	76.8%	authorship only

253 **The release.** The supplementary bundle holds the same-generator split (1,444 honest programs,
254 1,112 edit-arm and 1,112 independent-arm attack programs written by one model, plus every one
255 of the 11,565 draws of the independent arm’s retry pool), both eligibility rules as pair lists under an
256 outcome-blind selection (623 passing and 14 firing independent-arm pairs) and under the benchmark’s
257 own survivor selection (1,046 and 99, or 106 counting a crash on the trigger as firing), a dataset card,
258 and a manifest of row counts and checksums, under the MIT licence.

259 **The test harness and the pass flag.** Gate S requires executing generated code, so we built a runner
260 that feeds each shipped input case to a solution under per-case timeouts and address-space limits.
261 It agrees with the artifact’s own `solution_passes_tests` on 95.84% of the analysis pool and is
262 strictly conservative there — it never credits a solution the artifact rejects (0 such cases). We use *our*
263 *own* flag for both classes: a checker stricter on one class than the other would manufacture exactly
264 the artifact this paper documents. The headline does not depend on that choice — under the artifact’s
265 flag $B = 0.871$ at $n = 1,444$, under ours $B = 0.873$ at $n = 1,406$, within one standard deviation of




Figure 2: **The structural floor in five same-generator agentic settings.** Each setting’s content-blind classifier — six count features, no content read — scored *within task* against its own permutation null. No authorship confound is possible here: claude-sonnet-4, gemini-2.5-pro and gpt-5 write both classes in all five. bash sits on its null at 0.504; agentdojo reaches 0.802 from character count alone. **The task-group count beside each setting matters:** shade and iac average over five task groups each while bash averages over 247, so these floors are not equally well determined. We report the within-task column because our own pooled numbers partly broke on task composition, and we broke them ourselves. The permutation nulls shown were computed with a pooled estimator rather than the folded within-task statistic; they bound the chance level but are not a like-for-like null for each dot. These corpora declare no licence; we redistribute nothing and report aggregates only.

each other. Our release carries one MIT-licensed file of published monitor scores, with its notice; the monitor scores in Figure 1 are runs/x0/monitor-summary.json of that release.

The token test, run and failed. We attempted to confirm the token-level mechanism of [14] in our setting and it does not replicate: separation -0.013 — the wrong sign — at permutation $p = 0.956$. The reason is register. Their human class is general-purpose repository code, whereas *both* of our classes are competitive-programming code, so an identifier axis that separates general human code from model code does not separate these two populations. Only 8 of their 10 canonical tokens are in our vocabulary, so the test is underpowered; but the point estimate has the wrong sign, which is not what an underpowered true effect looks like. We report it because a reader who knows that work may run exactly this test.

A third channel, noted rather than developed. In the same five agentic settings a fifteen-word regular expression over the untrusted policy’s own scratchpad reaches 0.76–0.82, because prompted attack policies restate their side task in the reasoning the monitor is given. We note this only to bound the claim: the instrument belongs to [13], which validates it properly against an LLM judge and a length-matched null, and we defer to that treatment.