
What Survives the Filter? Trivial Baselines on Stored Honest/Attack Pairs in AI-Control Monitoring

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Trusted monitoring in AI control is evaluated on stored honest/attack pairs. The
2 field already knows these pairs leak: it filters to working solutions and strips
3 comments. We ask what that preprocessing leaves behind. On the field’s own
4 filtered, comment-stripped split, a bag-of-words classifier scores above roughly
5 four-fifths of the open-weight monitors ranked on it, and a classifier that reads no
6 code at all — only problem metadata — scores above most of them. We trace the
7 residue to authorship rather than sabotage: the honest class is human-written while
8 every attack is model-written. The signal transfers across backdoor generators, is
9 nearly as strong on inert backdoors as on working ones, and survives the field’s
10 comment stripper. Regenerating only the honest half does not isolate content:
11 with an open-weight model writing honest solutions against the stored backdoors,
12 separability *rises*. Holding authorship constant by construction, with one model
13 writing both classes, removes about a third of the separability and leaves the rest:
14 the same-generator control excludes authorship as the whole story and attributes
15 the remainder to surface content of the attack class. We recommend that stored-pair
16 evaluations report partial-input and bag-of-words floors against permutation nulls,
17 use same-generator honest samples, and drop inert backdoors from both classes.

18 1 Introduction

19 Trusted monitoring is the backbone protocol of AI control: a weaker trusted model reads an untrusted
20 model’s output and flags sabotage [4]. It is developed, tuned and ranked almost entirely on *stored*
21 honest/attack pairs — fixed corpora of solutions, some carrying a planted backdoor [12, 19, 9]. The
22 field already knows these pairs leak. Koran et al. keep only working honest solutions and working
23 backdoors and strip every comment, with the rationale stated [12]; Control Tax reports the pass-tests
24 pre-filter and appendices “Monitoring failed backdoors” [19]; [18] is an entire paper on the comment
25 channel. The sharpest treatment is Jha’s, whose released harness ships a `-leak-check` flag with two
26 trivial baselines already inside it and concludes that stripping is mandatory rather than cosmetic: a
27 #-count detector falls from 0.6957 to 0.4983 once comments go, and code length alone sits at 0.4726
28 [9]. This paper starts from that work.

29 What nobody has measured is what the preprocessing leaves behind. On the field’s own filtered,
30 comment-stripped split, a bag-of-words classifier reaches AUROC 0.806 — above 20 of the 24
31 open-weight monitors ranked on that same split, whose median is 0.609 — and a classifier that reads
32 no code at all, only problem difficulty, test count and question length, reaches 0.772, above 20.

33 Much of the residue is authorship, not sabotage. Honest solutions come from APPS [7] and are
34 human-written; every backdoor is model-written. The classifier transfers across backdoor generators
35 and is nearly as strong on *inert* backdoors as on working ones — neither is a property of a backdoor.
36 That bag-of-words separates human-written from model-written code is established out-of-field [14];

the known-ness of the mechanism is the argument, not a dent in it. The neighbouring field already applies the control missing here — CodeMirage gates its two classes on line count and character length [5], MultiAIGCD pairs to shared problem definitions [1] — so unmatched pairs are a defect by a published standard, not one we invented. Adjacent control work attacks other axes: [20] subverts monitors with injected text, [3] shows attacker-side selection inflates safety estimates, and [15] removes transcript length as a nuisance covariate, where we publish a content-blind AUROC as a per-benchmark floor.

We contribute (i) a protocol $\times$ baseline measurement of what survives the field’s preprocessing; (ii) a same-generator control attributing roughly a third of the separability to authorship and the rest to surface content of the attack class, with a negative result on the way: regenerating one half of a stored pair swaps the confound rather than removing it; (iii) a second selection artifact the filter installs by defining the negative class through attack failure; and (iv) a preprocessing checklist and a released split. **Scope:** this concerns monitor training, selection and benchmarking on *stored* pairs. ControlArena’s live evaluations generate both classes and are not implicated.

2 Setup

The confound is legible in the schema. The canonical stored-pair artifact for APPS-based control evaluations [8] records `model` and `source` on every backdoor and no generator field on any honest solution (Table 1, appendix): the honest class is taken from APPS [7] and names no model; every backdoor is model-written.

Artifacts and provenance. We use seven public corpora; their per-class provenance is the object of study: in the APPS family the honest class is human and the attack class is one of ten named models, while in the five agentic settings the same three models write *both* classes. **Six of the seven, from two publishers, declare no licence** [11, 17]; we report only aggregates over them and redistribute no corpus.

Evaluation surface, and naming protocols. Our headline protocol is the filtered, comment-stripped split of [12]: 1,209 train and 1,375 test items (172 attack, 1,203 honest), labelled by `backdoor_works`, reconstructed with the released code of [9]. It is the surface on which 24 open-weight monitors have published per-item scores. Five populations (the raw artifact, this split, pooled, matched pairs, within-task) and two metrics appear below; every comparison names its split, metric and preprocessing. Generated code is scored by our own runner, which agrees with the artifact’s `solution_passes_tests` on 95.84% of the analysis pool and is applied to both classes (appendix A).

3 Results

Table 2 (appendix) reports what a trivial classifier recovers under each protocol. The fix works: code length is driven to chance on the filtered split (0.4726) and the #-count detector falls from 0.6957 to 0.4983. What it does not reach, on that same split, is a bag-of-words floor of 0.806 (95% CI 0.7704–0.8409 by bootstrap over problems; permutation null 95th percentile 0.5465, $p < 0.005$), above 20 of the 24 monitors by point estimate and 18 of them with a paired-bootstrap interval excluding zero, and a no-code floor of 0.772 (CI 0.7319–0.8005; null 0.5545, $p < 0.005$), above 20 and 17 respectively (Figure 1): the preprocessing removed the carrier it targeted and not the two beside it.

The signal is authorship, not sabotage. Three measurements point the same way; every matched-protocol number here is comment-stripped on both classes, and the unstripped values, inflated by the comment channel, are in appendix A. First, the classifier *transfers across backdoor generators* at 0.969 pairwise: trained on one attacker’s backdoors it identifies another’s, so what it has learned is not any single model’s sabotage style. Second, it is nearly as strong on **inert** backdoors (0.879) as on working ones (0.935) — if it were detecting sabotage, backdoors that do not function should be much harder, and they are not. Third, the residual carrier is stylistic rather than semantic: `__main__` appears in 0.116 of attack samples against **0.000** of honest ones, replicated on the train split (152 versus 1 of 1,209) and untouched by the stripper. None of these is a property of a backdoor. A fourth,

token-level test [14] does not replicate here (separation -0.013 , permutation $p = 0.956$; appendix A), a register effect; and one authorship leak, author bylines inside string literals, survives the field’s own comment stripper (appendix A).

Regenerating one half of a pair does not isolate content. Under the identical pipeline the human-honest baseline is $B = 0.871$ over $n = 1,444$ pairs, and the routing rule and its bands were fixed before any arm ran: retained signal $\rho = (G - 0.5)/(B - 0.5)$, with $\rho \leq 0.33$ reading as authorship and $\rho > 0.72$ as content. With an open-weight model writing the honest class against the same backdoors ($n = 668$, $B' = 0.874$ on those problems), separation *rises* to $G = 0.962$, $\rho = 1.234$; TF-IDF gains $+0.088$ while the length carrier collapses to 0.489 . Substituting the honest author replaced one authorship confound with a larger one.

Holding authorship constant needs care about what else changes. We generate the attack class twice from one honest class: an **edit** arm shown its own honest solution, whose pairs are near-copies, and an **independent** arm shown only the problem statement, whose pairs match the stored regime of independent programs (Table 3, appendix). Only the second changes authorship alone.

Holding authorship constant removes about a third of the separability and leaves the rest. With one model writing both classes — one outcome-blind draw per problem, the attack class defined, as in the benchmark, by intent plus passing the public tests, and B' inert-inclusive over the same problems — $S = 0.751$ (95% CI $0.7295\text{--}0.7713$; $n = 623$ problem-disjoint pairs) against $B' = 0.8658$, so $\rho = 0.687$ ($0.625\text{--}0.755$). The interval excludes the collapse band and spans the content line, so the pre-registered rule returns *inconclusive between partial and content*. Authorship accounts for roughly a third of the separability, and what survives is lexical: length-only falls to 0.5851 while TF-IDF holds. It is not the trigger: masking numerals and comparison operators leaves S unchanged, while normalising identifiers takes it to 0.700 ($\rho = 0.54$) with B' untouched, and the top attack-side features are `def main` and `data` (67% vs 39% of pair sides) — the layout the attack prompt elicits. The arm holds the model constant, not the prompt, so ρ is an upper bound on content. The pairs are independent programs (similarity 0.659 ; stored pairs on the same problems 0.606), and at $n = 623$ B under the identical pipeline is 0.8521 ± 0.0055 , so the drop is not sample size. The strict row — attacks that pass *and* fire, kept by the benchmark’s own retry-and-keep procedure, the filtering this paper criticises, inherited deliberately — agrees: $\rho = 0.921$ ($0.662\text{--}1.089$) from $S = 0.7652$ ($0.7016\text{--}0.8000$; $n = 99$) against its own $B' = 0.7877$, read through ρ alone because B at that n is 0.7807 ± 0.0248 . Retention is not the artifact: survivors and non-survivors separate at only 0.62 , and S with non-firing attempts on the same problems is 0.7595 against 0.7687 . The edit arm bounds S under near-copies (similarity 0.928): 0.6266 , $\rho = 0.371$. A trusted-monitor prompt to the same model — the one that wrote both classes — scores 0.818 ($0.7630\text{--}0.8686$) on the strict pairs, within the baseline’s interval, but 0.566 ($0.5507\text{--}0.5833$) on the headline set, where TF-IDF holds 0.751 ; the two are independent pair by pair ($\phi = 0.04$) and read different things — the layout, and the backdoor. Our single-shot attacks fire given a pass at 2.5% against the stored attacker’s 67.2% , which executes and iterates: the stored attack class is a survivor set. Intervals are bootstraps over problems, resampling S and B' jointly for ρ ; this replaced a ± 2 fold-sd rule before the retry pool was read.

The filter installs a second artifact. Labelling by `backdoor_works` defines the negative class by *attack failure* rather than by honesty: 80.7% of the test split’s labelled backdoors do not work (32.8% of those that pass the public tests) and are relabelled honest, so most of the “honest” class consists of failed attacks. Attack problems carry roughly half as many test cases as honest ones (20.95 versus 38.85), consistent with a problem with fewer tests being easier to backdoor undetectably, so difficulty and test count are correlated with the label before any code is read — which is where the content-blind floor of 0.772 comes from.

The negative control comes free. In five same-generator agentic settings the same three models write both classes, so no authorship confound is possible and any surviving floor is structural (Figure 2, appendix). `bash` sits on its permutation null at 0.504 while `agentdojo` reaches 0.802 from character count alone. A per-benchmark floor therefore separates corpora where the label is trivially recoverable from those where it is not.




Figure 1: **Two floors against the 24 open-weight monitors ranked on the same split.** All values are AUROC on Koran et al.’s 1,375-item test split (172 attack / 1,203 honest), comments stripped. Monitor scores from [9] (MIT); median 0.609, max 0.937. **Two caveats.** (i) The 24 monitors are zero-shot prompted while the bag-of-words classifier is *trained* on the benchmark’s own paired train split — this is a statement about what the corpus makes recoverable, in the sense of [6], not a claim that bag-of-words is a better monitor. (ii) The no-code floor is fitted *within* the test split by 5×5 cross-validation, because the released train/val split is paired and therefore metadata-balanced, so its permutation null shuffles labels for fitting and scoring alike; on a single held-out slice of the same items it is 0.716 (CI 0.6476–0.7796; $p = 0.030$ against its own null), above 18 of 24, 12 with a paired interval excluding zero.

4 Recommendations

Following Kapoor and Narayanan [10], who treat leakage as a checklist rather than a verdict:

1. **Generate both classes with the same model.** If the attack class is model-written, the honest class must be too; as our arm-1 null shows, regenerating only one side substitutes one authorship gap for another. This is precedent [5, 1], not invention.
2. **Report a partial-input floor and a bag-of-words floor beside every monitor number,** each against a permutation null.
3. **Do not define the negative class by attack failure.** Drop inert backdoors from both classes rather than relabelling a failed attack as honest.
4. **State the split, the metric and the preprocessing** whenever monitors are compared, and **release the split and the scripts.** Ours accompany this paper as supplementary material under MIT: the same-generator split (1,444 honest, 1,112 edit-arm and 1,112 independent-arm programs from one model, plus all 11,565 retry-pool draws), pair lists under both selection rules — outcome-blind (623 passing, 14 firing independent-arm pairs) and the benchmark’s own survivor selection (1,046 and 99) — a dataset card and a manifest (appendix A).

5 Limitations

A high floor is evidence about a corpus, not a verdict on results computed on it. Feng et al. [2] warn that the *failure* of a partial-input baseline does not show a dataset is artifact-free; the converse binds us: our floors show the label is recoverable without reading a backdoor, in the sense of Gururangan et al. [6] and Poliak et al. [16], not that any monitor exploits that.

The same-generator arm is a later model behind a paid API. The stored attacker, claude-3-7-sonnet-thinking, is retired (404 against a live key), so both same-generator classes come from a later model of the same lineage, called through a metered commercial API. That API does not expose its sampling parameters, so the arm is not bit-reproducible as the open-weight arms are, and thinking was off although the stored attacker was a thinking model. We do not length-match: in deployment the untrusted model chooses its trajectory length.

Scope, provenance and the harness. Whether stored-pair corpora leak this way in general is not established here. Five of our seven corpora are one user’s repackaging; claims are scoped to those releases. Two interpreter-era artifacts fall on opposite classes (a `stdin` without `.buffer`, fixed; `fractions.gcd`, removed in Python 3.9, not fixable); the same-generator arm is immune, *B* is not.

References

- [1] Basak Demirok, Mucahid Kutlu, and Selin Mergen. MultiAIGCD: A comprehensive dataset for AI generated code detection covering multiple languages, models, prompts, and scenarios. 2025. arXiv:2507.21693v1.
- [2] Shi Feng, Eric Wallace, and Jordan Boyd-Graber. Misleading failures of partial-input baselines. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2019.
- [3] Catherine Ge-Wang, Tyler Crosse, Benjamin Hadad, Joachim Schaeffer, Ram Potham, and Tyler Tracy. Attack selection in agentic AI control evaluations meaningfully decreases safety. 2026. arXiv:2606.06529.
- [4] Ryan Greenblatt, Buck Shlegeris, Kshitij Sachan, and Fabien Roger. AI control: Improving safety despite intentional subversion. 2023. arXiv:2312.06942.
- [5] Hanxi Guo, Siyuan Cheng, Kaiyuan Zhang, Guangyu Shen, and Xiangyu Zhang. CodeMirage: A multi-lingual benchmark for detecting AI-generated and paraphrased source code from production-level LLMs. 2025. arXiv:2506.11059v1. No venue stated.
- [6] Suchin Gururangan, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel R. Bowman, and Noah A. Smith. Annotation artifacts in natural language inference data. In *Proceedings of NAACL-HLT*, 2018.
- [7] Dan Hendrycks, Steven Basart, Saurav Kadavath, Mantas Mazeika, Akul Arora, Ethan Guo, Collin Burns, Samir Puranik, Horace He, Dawn Song, and Jacob Steinhardt. Measuring coding challenge competence with APPS. In *NeurIPS Datasets and Benchmarks Track*, 2021. arXiv:2105.09938. The upstream source of the human-written honest class, distributed as the HuggingFace dataset `codeparrot/apps`.
- [8] Rogan Inglis. `apps-control-arena`. HuggingFace dataset `RoganInglis/apps-control-arena`. MIT licence. 5,000 rows; honest class from `codeparrot/apps`., 2025.
- [9] Anik Jha. Decorrelation is not complementarity: Skill, not lineage, governs trusted-monitor ensembles. 2026. arXiv:2608.16190. Code: <https://github.com/anik-jha/challenger-panels> (MIT).
- [10] Sayash Kapoor and Arvind Narayanan. Leakage and the reproducibility crisis in ML-based science. 2022. arXiv:2207.07048.
- [11] Eugene Koran. `apps-control-arena-high-quality`. HuggingFace dataset `eugenekoran/apps-control-arena-high-quality`. **No licence declared.** 1,209 train / 1,375 test., 2026.
- [12] Eugene Koran, Yejun Yun, Samantha Tetef, Benjamin Arnav, and Pablo Bernabeu-Pérez. Ensemble monitoring for AI control: Diverse signals outweigh more compute. 2026. arXiv:2605.15377v2.
- [13] Bryce Little. Length penalties make chain-of-thought less monitorable. 2026. arXiv:2607.09786v3.
- [14] Daniil Orel, Dilshod Azizov, Indraneil Paul, Yuxia Wang, Iryna Gurevych, and Preslav Nakov. AICD bench: A challenging benchmark for AI-generated code detection. 2026. arXiv:2602.02079v1. No venue stated on the abs page.
- [15] Lucas Pinto. Calibration-family overfit: Why trusted sabotage monitors don’t transfer across lineages. In *COLM 2026 Workshop on AI Measurement Science*, 2026. arXiv:2607.06596v3.

- [16] Adam Poliak, Jason Naradowsky, Aparajita Haldar, Rachel Rudinger, and Benjamin Van Durme. Hypothesis only baselines in natural language inference. In *Proceedings of the Seventh Joint Conference on Lexical and Computational Semantics (*SEM)*, 2018.
- [17] Aditya Sinha. control_arena_{bash,shade,iac,rogue_eval,agentdojo}. HuggingFace datasets adityaasinha28/control_arena_*. **No licence declared**. 1,500 rows each., 2026.
- [18] Ashwin Sreevatsa, Sebastian Prasanna, and Cody Rushing. Basic legibility protocols improve trusted monitoring. 2026. arXiv:2602.10153.
- [19] Mikhail Terekhov, Zhen Ning David Liu, Caglar Gulcehre, and Samuel Albanie. Control tax: The price of keeping AI in check. 2025. arXiv:2506.05296v3.
- [20] Mikhail Terekhov, Alexander Panfilov, Daniil Dzenhaliou, Caglar Gulcehre, Maksym Andriushchenko, Ameya Prabhu, and Jonas Geiping. Adaptive attacks on trusted monitors subvert AI control protocols. In *International Conference on Learning Representations (ICLR)*, 2026. arXiv:2510.09462v2.

A Supplementary tables and figures

The schema. Table 1 is the record schema of RoganInglis/apps-control-arena [8]. The attack class carries model and source, naming the system that produced each backdoor; the honest class carries no such field, because there is no generator to name. The two APPS-backdoor corpora we use are two views of this one artifact: every problem in the filtered release [11] is in [8], byte-identical except a re-validated backdoor_works. Among the agentic corpora we exclude AgentHarm, whose author is an invited speaker at this venue.

Table 1: The confound is legible in the dataset schema, with no measurement at all. The attack class records which model wrote it; the honest class has no generator field. Verified against the shard schema of RoganInglis/apps-control-arena.

struct	sub-fields
solutions	code, compiles, passes_tests
backdoors	backdoor_code, backdoor_input, backdoor_works, compiles, model, passes_tests, source

The protocol $\times$ baseline table. Table 2 is the full measurement behind §3. Rows are not comparable across protocols; the preprocessing column says whether comments were stripped before the baseline read the code. Pairwise accuracies are given to three decimals because their resolution over 2,641 pairs is one pair. Intervals in §3 are bootstraps over problems (2,000 draws) and nulls are 200 label permutations; the held-out no-code row is barely above its own null (0.716 against a 95th percentile of 0.7023, $p = 0.030$), so the no-code claim rests on the cross-validated protocol.

Comment stripping on the matched-protocol numbers. The matched-pair and transfer measurements were first made with comments in. Every value in the authorship paragraph of §3 was re-measured with the same stripper applied to both classes, over five problem-disjoint splits (each transfer row is one fixed split). With comments kept the same quantities were transfer 0.994, working 0.979, inert 0.955, all pairs 0.962 and surface-feature AUROC 0.8306; stripped they are 0.969, 0.935, 0.879, 0.897 and 0.7259. Every one was inflated by the comment channel, which is this paper’s thesis applied to itself. The #-presence marker, 0.685 against 0.176 before stripping, closes to 0.007 against 0.008 after and is dropped; __main__ is unchanged. One authorship leak survives the stripper because string literals are not comments: in the comment-stripped train split an author byline (Author :) survives in 23 of 1,209 honest samples against 2 attack samples, and a triple-quoted block in 43 against 9 — too rare to drive 0.806, but a literal record of authorship. Stripped, the matched-protocol TF-IDF AUROC is 0.8560, beside $B = 0.871$ and the Koran-split 0.806, all three on stripped code. The other cross-generator transfer rows, stripped, are 0.866 (o3-mini to claude-3-7-sonnet, 82 test pairs), 0.844 (the reverse) and 0.764 (o1-mini to gemini-2.0-flash).

Table 2: What the field’s preprocessing leaves behind. Every row is measured under the protocol named in column one; rows are not comparable across protocols. *Prep.* says whether comments were stripped before the baseline read the code; *no code* means the baseline never reads the program. Pairwise accuracies are given to three decimals, their resolution over the pair count.

Protocol	Baseline	Value	Prep.	Against
<i>AUROC</i>				
Koran’s exact split	TF-IDF word 1–2gram	0.8063	stripped	above 20 of 24
same	TF-IDF word 1–2gram	0.8914	kept	above 22 of 24
same, reads no code	difficulty + #tests + $ q $	0.7715 ± 0.0389	no code	above 20 of 24
same, held-out slice	difficulty + #tests + $ q $	0.7163	no code	above 18 of 24
same	code length only	0.4726	stripped	chance
same	#-count detector	0.4983	stripped	chance
same	#-count detector	0.6957	kept	—
Raw artifact, problem level	all metadata	0.8917	no code	—
same	difficulty only	0.6823	no code	—
same	<code>solution_passes_tests</code>	0.7881	no code	—
Matched pairs (P3)	all metadata	0.5000	no code	exactly chance by construction
<i>Pairwise accuracy</i>				
DSPy matched, problem-disjoint	TF-IDF word 1–2gram	0.897	stripped	tutorial monitor 0.69–0.81
working backdoors only	TF-IDF word 1–2gram	0.935	stripped	—
inert backdoors only	TF-IDF word 1–2gram	0.879	stripped	the signal is not sabotage
transfer, Control Tax → GPT-4-era set	TF-IDF word 1–2gram	0.969	stripped	—
all pairs, comments kept	TF-IDF word 1–2gram	0.962	kept	inflated by the comment channel
working, comments kept	TF-IDF word 1–2gram	0.979	kept	inflated by the comment channel
inert, comments kept	TF-IDF word 1–2gram	0.955	kept	inflated by the comment channel
transfer, comments kept	TF-IDF word 1–2gram	0.994	kept	inflated by the comment channel
same, raw artifact	code length only	0.826	kept	—

253 **Pair similarity, and why the same-generator control needs two arms.** An attack class built
254 by editing its own honest solution changes pair similarity as well as authorship (Table 3); the
255 independent arm sits in the stored benchmark’s regime and if anything slightly further from it, so it
256 cannot manufacture a low S by making the classes alike.

Table 3: Character similarity of honest/attack pairs, comments stripped. The edit arm’s classes are near-copies; the independent arm reproduces the stored benchmark’s regime, where the two classes are largely independent programs.

Attack class	mean sim.	> 0.95	< 0.60	changes vs. the benchmark
stored benchmark	0.606	3.4%	45.4%	— (the regime to match)
edit (shown its own solution)	0.928	64.2%	4.0%	authorship <i>and</i> similarity
independent (problem only)	0.659	3.4%	34.8%	authorship only

257 **The five same-generator agentic settings.** Figure 2 shows each setting’s content-blind floor against
258 its own permutation null.

259 **The release.** The supplementary bundle holds the same-generator split (1,444 honest programs,
260 1,112 edit-arm and 1,112 independent-arm attack programs written by one model, plus every one
261 of the 11,565 draws of the independent arm’s retry pool), both eligibility rules as pair lists under an
262 outcome-blind selection (623 passing and 14 firing independent-arm pairs) and under the benchmark’s
263 own survivor selection (1,046 and 99, or 106 counting a crash on the trigger as firing), a dataset card,
264 and a manifest of row counts and checksums, under the MIT licence.




Figure 2: **The structural floor in five same-generator agentic settings.** Each setting’s content-blind classifier — six count features, no content read — scored *within task* against its own permutation null. No authorship confound is possible here: claude-sonnet-4, gemini-2.5-pro and gpt-5 write both classes in all five. bash sits on its null at 0.504; agentdojo reaches 0.802 from character count alone. **The task-group count beside each setting matters:** shade and iac average over five task groups each while bash averages over 247, so these floors are not equally well determined. We report the within-task column because our own pooled numbers partly broke on task composition, and we broke them ourselves. The permutation nulls shown were computed with a pooled estimator rather than the folded within-task statistic; they bound the chance level but are not a like-for-like null for each dot. These corpora declare no licence; we redistribute nothing and report aggregates only.

265 **The test harness and the pass flag.** Gate S requires executing generated code, so we built a runner
 266 that feeds each shipped input case to a solution under per-case timeouts and address-space limits.
 267 It agrees with the artifact’s own `solution_passes_tests` on 95.84% of the analysis pool and is
 268 strictly conservative there — it never credits a solution the artifact rejects (0 such cases). We use *our*
 269 *own* flag for both classes: a checker stricter on one class than the other would manufacture exactly
 270 the artifact this paper documents. The headline does not depend on that choice — under the artifact’s
 271 flag $B = 0.871$ at $n = 1,444$, under ours $B = 0.873$ at $n = 1,406$, within one standard deviation of
 272 each other. Our release carries one MIT-licensed file of published monitor scores, with its notice; the
 273 monitor scores in Figure 1 are runs/`x0/monitor-summary.json` of that release.

274 **The token test, run and failed.** We attempted to confirm the token-level mechanism of [14] in our
 275 setting and it does not replicate: separation -0.013 — the wrong sign — at permutation $p = 0.956$.
 276 The reason is register. Their human class is general-purpose repository code, whereas *both* of our
 277 classes are competitive-programming code, so an identifier axis that separates general human code
 278 from model code does not separate these two populations. Only 8 of their 10 canonical tokens are in
 279 our vocabulary, so the test is underpowered; but the point estimate has the wrong sign, which is not
 280 what an underpowered true effect looks like. We report it because a reader who knows that work may
 281 run exactly this test.

282 **A third channel, noted rather than developed.** In the same five agentic settings a fifteen-word
 283 regular expression over the untrusted policy’s own scratchpad reaches 0.76–0.82, because prompted
 284 attack policies restate their side task in the reasoning the monitor is given. We note this only to bound
 285 the claim: the instrument belongs to [13], which validates it properly against an LLM judge and a
 286 length-matched null, and we defer to that treatment.